

```

import sys
import numpy as np
import pandas as pd
import openpyxl
import matplotlib.pyplot as plt
from pathlib import Path

# === Define Versions ===
print(sys.executable)
print("pandas:", pd.__version__, "openpyxl:", openpyxl.__version__)

# === Load dataset ===
candidate = Path(r"C:\Users\jaishva\Downloads\Teamwork-regress w 2 losses (dataset)-1.xlsx")
df = pd.read_excel(candidate, engine="openpyxl")
print(df.head())

# === Preparing data (y = w*x + b) ===
X = np.c_[df["x"].to_numpy(), np.ones(len(df))]
y = df["y"].to_numpy().reshape(-1,1)

# === Gradient Descent Implementation ===
def gd_linear_regression(X, y, loss="mse", lr=0.01, delta=1.0, max_iters=1000,
                        tol_param=1e-6, tol_loss=1e-8):
    """
    Gradient descent for linear regression.
    Supports MSE and Huber losses.
    """
    n, d = X.shape
    theta = np.zeros((d,1)) # init params
    losses, param_deltas = [], []
    k_param = k_loss = None

    for k in range(max_iters):
        y_pred = X @ theta
        error = y_pred - y

        # --- Loss + gradient ---
        if loss == "mse":
            L = np.mean(error**2)
            grad = (2/n) * (X.T @ error)
        elif loss == "huber":
            abs_err = np.abs(error)
            quadratic = abs_err <= delta
            L = np.mean(
                np.where(quadratic,
                        0.5*error**2,
                        delta*(abs_err - 0.5*delta))

```

```

        )
        grad = (X.T @ np.where(quadratic, error, delta*np.sign(error))) /
n
    else:
        raise ValueError("loss must be 'mse' or 'huber'")

    losses.append(L)

    # --- Updates ---
    theta_new = theta - lr*grad
    param_delta = np.linalg.norm(theta_new - theta)
    param_deltas.append(param_delta)
    theta = theta_new

    # --- Stopping criteria ---
    if k_param is None and param_delta < tol_param:
        k_param = k
    if k_loss is None and k > 0 and abs(losses[-2] - losses[-1]) <
tol_loss:
        k_loss = k
    if k_param is not None or k_loss is not None:
        break

    return theta, losses, param_deltas, k_param, k_loss

# == Let us run our experiments ==
lrs = [0.01, 0.05, 0.2]
delta = 0.5

for loss in ["mse", "huber"]:
    print(f"\n=== {loss.upper()} ===")
    plt.figure()
    for lr in lrs:
        theta, losses, param_deltas, k_param, k_loss = gd_linear_regression(
            X, y, loss=loss, lr=lr, delta=delta, max_iters=3000,
            tol_param=1e-6, tol_loss=1e-8
        )
        print(f"lr={lr:<4}  theta=[{theta[0,0]:.4f}, {theta[1,0]:.4f}]  "
              f"final_loss={losses[-1]:.6f}  "
              f"stop(param)={k_param}  stop(loss)={k_loss}  iters={len(losses)}")

    # Visualize all LR's on one figure
    it = np.arange(len(losses))
    plt.plot(it, losses, label=f"lr={lr}")
    if k_param is not None and k_param < len(losses):
        plt.scatter([k_param], [losses[k_param]], marker='o')
    if k_loss is not None and k_loss < len(losses):

```

```

plt.scatter([k_loss], [losses[k_loss]], marker='x')

plt.xlabel("Iteration")
plt.ylabel("Training loss")
plt.title(f"Learning curves - {loss.upper()} (o: param tol, x: loss tol)")
plt.legend()
plt.grid(True, alpha=0.3)
plt.show()

# === Let us now compare the values with Least Squares solution ===
theta_ls, *_ = np.linalg.lstsq(X, y, rcond=None)
print("LS theta:", theta_ls.ravel())

theta_mse, *_ = gd_linear_regression(X, y, loss="mse", lr=0.05,
max_iters=3000)
print("GD theta (MSE, lr=0.05):", theta_mse.ravel())

print("||GD - LS||:", np.linalg.norm(theta_mse - theta_ls))

```

## Terminal Output

C:\Users\jaishva\Assignment\venv\Scripts\python.exe

pandas: 2.3.3 openpyxl: 3.1.5

x y

0 0.417411 0.841049

1 0.222108 0.556829

2 0.119865 0.518283

3 0.337615 0.788053

4 0.942910 1.067603

=== MSE ===

lr=0.01 theta=[0.7266, 0.4607] final\_loss=0.011105 stop(param)=None stop(loss)=None  
iters=3000

lr=0.05 theta=[0.7305, 0.4589] final\_loss=0.011102 stop(param)=None stop(loss)=721 iters=722

lr=0.2 theta=[0.7320, 0.4581] final\_loss=0.011102 stop(param)=None stop(loss)=204 iters=205

=== HUBER ===

lr=0.01 theta=[0.6768, 0.4844] final\_loss=0.005686 stop(param)=None stop(loss)=None  
iters=3000

lr=0.05 theta=[0.7276, 0.4603] final\_loss=0.005552 stop(param)=None stop(loss)=1259 iters=1260

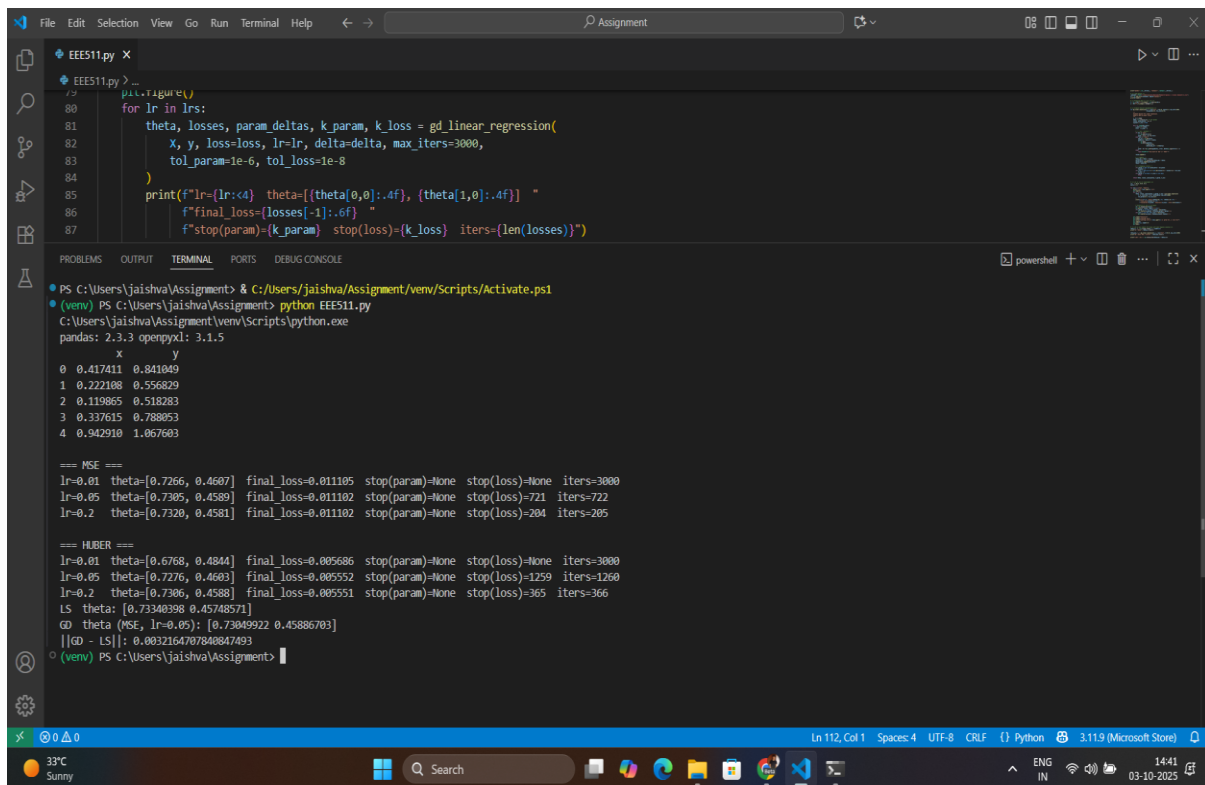
lr=0.2 theta=[0.7306, 0.4588] final\_loss=0.005551 stop(param)=None stop(loss)=365 iters=366

LS theta: [0.73340398 0.45748571]

GD theta (MSE, lr=0.05): [0.73049922 0.45886703]

||GD - LS||: 0.0032164707840847493

Proof:-



The screenshot shows a VS Code editor with a file named `EEES11.py` open. The script defines a function `gd_linear_regression` and iterates over learning rates `lr` in `[0.01, 0.05, 0.2]`. The output in the terminal shows the results for each learning rate, including the final loss, the number of iterations, and the theta values. The output also shows the LS theta values and the difference between GD and LS theta values.

```
EEES11.py
def gd_linear_regression(X, y, loss=loss, lr=lr, delta=delta, max_iters=3000, tol_param=1e-6, tol_loss=1e-8):
    theta, losses, param_deltas, k_param, k_loss = gd_linear_regression(X, y, loss=loss, lr=lr, delta=delta, max_iters=3000, tol_param=1e-6, tol_loss=1e-8)
    print(f"lr={lr} theta=[{theta[0]:.4f}, {theta[1]:.4f}] final_loss={losses[-1]:.6f} stop(param)={k_param} stop(loss)={k_loss} iters={len(losses)}")

if __name__ == "__main__":
    for lr in lrs:
        theta, losses, param_deltas, k_param, k_loss = gd_linear_regression(X, y, loss=loss, lr=lr, delta=delta, max_iters=3000, tol_param=1e-6, tol_loss=1e-8)
        print(f"lr={lr} theta=[{theta[0]:.4f}, {theta[1]:.4f}] final_loss={losses[-1]:.6f} stop(param)={k_param} stop(loss)={k_loss} iters={len(losses)}")

# Output:
PS C:\Users\jaishva\Assignment> & C:\Users\jaishva\Assignment\venv\Scripts\Activate.ps1
(venv) PS C:\Users\jaishva\Assignment> python EEES11.py
C:\Users\jaishva\Assignment\venv\Scripts\python.exe
pandas: 2.3.3 openpyxl: 3.1.5

x      y
0  0.417411  0.841049
1  0.222108  0.256820
2  0.119865  0.518283
3  0.337615  0.788853
4  0.942910  1.067603

=== MSE ===
lr=0.01 theta=[0.7266, 0.4607] final_loss=0.011105 stop(param)=None stop(loss)=None iters=3000
lr=0.05 theta=[0.7306, 0.4588] final_loss=0.011102 stop(param)=None stop(loss)=721 iters=722
lr=0.2 theta=[0.7320, 0.4581] final_loss=0.011102 stop(param)=None stop(loss)=284 iters=285

=== L2R ===
lr=0.01 theta=[0.6768, 0.4844] final_loss=0.005686 stop(param)=None stop(loss)=None iters=3000
lr=0.05 theta=[0.7276, 0.4603] final_loss=0.005552 stop(param)=None stop(loss)=1259 iters=1260
lr=0.2 theta=[0.7306, 0.4588] final_loss=0.005551 stop(param)=None stop(loss)=365 iters=366
LS theta: [0.73340398 0.45748571]
GD theta (MSE, lr=0.05): [0.73049922 0.45886703]
||GD - LS||: 0.0032164707840847493
(venv) PS C:\Users\jaishva\Assignment>
```

